

WebAuthn PRF-Based Vault Key Wrapping and Zero-Knowledge PWA Storage Architecture

Ian Pinto¹

Independent Researcher, Mumbai, India
ianpinto1980@gmail.com

Abstract. This paper presents a client-side security architecture that combines the Web Authentication (WebAuthn) pseudo-random function (PRF) extension with symmetric cryptography and carefully engineered browser storage patterns to achieve strong zero-knowledge properties for web-based credential vaults. The design is grounded in the TrustVault-PWA project and informs the reusable `webauthn-prf-zktv` library.

Keywords: WebAuthn, PRF, passkeys, zero-knowledge, IndexedDB, PWA, cryptography

1 Introduction

The Web Authentication (WebAuthn) standard and FIDO2 ecosystem provide strong public-key based authentication and passkeys that are increasingly used to replace passwords on the web. Recent extensions such as the pseudo-random function (PRF) allow a relying party to request a high-entropy symmetric secret derived from a passkey credential during an authentication ceremony, enabling end-to-end encryption of user data.[15,3]

Password managers and credential vaults seek to offer zero-knowledge security: server-side or client-side storage alone should never be sufficient to reveal user secrets. Achieving this property in a web progressive web application (PWA) requires careful coordination of cryptography, browser APIs, and local storage. This paper formalizes a PRF-based vault key wrapping scheme and an IndexedDB storage architecture designed toward this goal.

1.1 Contributions

The main contributions are:

- A formal description of a WebAuthn PRF-based vault key wrapping scheme suitable for PWAs.
- An IndexedDB schema and migration strategy that reinforce zero-knowledge properties.
- A reusable, open-source implementation, the `webauthn-prf-zktv` library, that packages these patterns.[12]
- Discussion of post-quantum considerations and migration paths.

2 Background and Related Work

2.1 WebAuthn, Passkeys, and PRF

WebAuthn [14] is the W3C standard, developed jointly with the FIDO Alliance’s CTAP2 protocol, for challenge–response authentication with asymmetric credentials held by an *authenticator* (a platform biometric sensor, a security key, or a synced passkey provider). *Passkeys* are discoverable WebAuthn credentials, typically synchronized across a user’s devices by the platform vendor. The **prf** client extension [14,15] exposes the CTAP2 **hmac-secret** capability to web applications: during a ceremony the authenticator evaluates a pseudo-random function (HMAC-SHA-256 over a credential-scoped key) at one or two caller-chosen 32-byte salts, and returns 32-byte outputs. Because evaluation requires the physical authenticator and (with **userVerification: required**) a user-verification gesture, the PRF output is a high-entropy symmetric secret that exists only transiently in client memory — a natural root for client-side encryption keys. Developer-facing treatments of the extension are given by Yubico [15] and the SimpleWebAuthn project [8].

2.2 PRF-Based Vault Encryption in Existing Systems

Bitwarden uses the PRF extension to unlock its vault without the master password: the PRF output decrypts a wrapped copy of the account encryption key [3]. The scheme in this paper differs in three ways: (i) the wrap key is derived through HKDF with an explicit versioned domain-separation label rather than used directly; (ii) client-side replay verification (challenge, origin, signature counter) is an integral step of the unlock algorithm; and (iii) the record format and storage schema are specified together, so that the zero-knowledge property can be argued about the *entire* persisted state. Generic PRF helper libraries and file-encryption demonstrations [1] integrate the extension but do not prescribe a complete vault architecture with fallback, migration, and storage semantics.

2.3 Password-Manager and Vault Architectures

Web credential vaults have explored several trust models orthogonal to this work’s. Horcrux [7] secret-shares credentials across multiple keystores so that no single server compromise reveals passwords; AuthStore [16] strengthens password-derived keys with a PAKE against untrusted storage servers; and SafeKeeper [6] anchors password confidentiality in a server-side trusted execution environment. These designs distribute or harden *server-side* trust, whereas the present scheme removes the server from the confidentiality boundary entirely: the unlock secret never leaves the authenticator, and the persisted client state is the only attack surface. Multi-factor derivation frameworks — MFDPG [10] and the MFKDF2 family [13], the latter of which explicitly contemplates the WebAuthn PRF as a derivation factor — are complementary: the wrapping scheme presented here can serve as the PRF-factor envelope inside such constructions

without changing the record format. Finally, Fábrega et al. [4] demonstrate that end-to-end encrypted password managers leak through non-cryptographic system surfaces (telemetry, caching, and pre-encryption compression or deduplication) under adversarial data injection. Those channels are out of scope for the storage layer formalized here, but Sect. 4.3 is deliberately scoped to *persisted state* for exactly this reason: applications adopting the library must still avoid pre-encryption compression, unpadded record sizes correlated to secrets, and metadata-bearing telemetry.

2.4 IndexedDB and Client-Side Storage

IndexedDB is the standard transactional object store available to web applications and service workers, and the only browser storage suitable for structured encrypted records of meaningful size. Data in IndexedDB is protected by the same-origin policy but is plaintext-at-rest from the application’s perspective: any principal that can read the origin’s storage (local malware, forensic imaging, a malicious extension with storage access) reads the stored bytes. The TrustVault-PWA project [11], from which this work is extracted, stores both secrets and their metadata encrypted in Dexie-backed IndexedDB and documents its residual plaintext surfaces, such as breach-detection (HIBP) hash prefixes retained for offline checking. The library presented here inherits that design posture: the storage layer never persists anything from which a key could be recomputed.

3 Threat Model

We consider four attacker capabilities:

- **Stored-data compromise.** The attacker obtains a full copy of persisted state — an IndexedDB dump, a device backup, or a synchronized replica — but does not control a live session.
- **Browser compromise.** The attacker executes script in the origin (XSS, malicious dependency) while the user is present.
- **Authenticator compromise.** The attacker gains use of the passkey (stolen unlocked device, cloned credential).
- **Harvest-now, decrypt-later.** The attacker records ciphertexts today in the expectation of future cryptanalytic (including quantum) capability.

The corresponding security goals are: *zero-knowledge storage* — stored-data compromise alone yields no secrets (Sect. 4.3); *bounded exposure* under browser compromise; *best-effort cloned-credential detection*; and *PQ-adequate* symmetric primitives at the storage layer (Sect. 7). Two of these goals deserve sharper statements than they usually receive. *Bounded exposure* is a real but coarse bound: non-extractable wrap keys and zeroized PRF outputs prevent a malicious in-origin script from exfiltrating key material for later offline use, but nothing prevents such a script from *using* an unlocked key — during its window of control it can decrypt every record the user unlocks, and applications should treat

origin compromise as full compromise of the unlocked session, mitigable only by compartmentalization (short unlock windows, per-record unlock ceremonies for high-value secrets). *Cloned-credential detection* via signature counters is likewise best-effort: synced passkeys commonly report a constant counter of zero, which disables the increase check entirely (Algorithm 2 skips it when $c'=0$). The library therefore exposes the authenticator’s backup-eligibility (BE) and backup-state (BS) flags via `readAuthenticatorFlags`, letting applications distinguish device-bound credentials — where a non-increasing counter is a strong clone signal — from synced passkeys, where it is vacuous; per-device continuity via the WebAuthn Level 3 `devicePubKey` extension is a natural upgrade once client support matures. Authenticator compromise combined with user verification defeats any client-side scheme and is out of scope.

4 PRF-Based Vault Key Wrapping

4.1 High-Level Design

Enrollment (`enrollVault`) and unlock (`unlockVault`) use the 32-byte WebAuthn PRF output as HKDF input keying material to derive a non-extractable AES-256-GCM *wrap key* K_w , which encrypts the 256-bit vault key K_v . The only persisted artifact is a `WrappedSecretRecord`:

$$R = (\text{scheme}, \text{ciphertext}, \text{nonce}, \text{salt}, [\text{kdfParams}])$$

where $\text{scheme} \in \{\text{prf-v1}, \text{pw-v1}\}$, the nonce is a fresh 96-bit random AES-GCM IV, and the salt is the PRF salt (`prf-v1`) or the scrypt salt (`pw-v1`). `kdfParams` is present iff $\text{scheme} = \text{pw-v1}$; the parser (`parseRecord`) rejects `prf-v1` records carrying it. PRF outputs and wrap keys are never stored: wrap keys are created with `extractable: false` in Web Crypto, and transient PRF-output buffers are zeroized (`.fill(0)`) in `finally` blocks.

HKDF uses SHA-256 with an *empty* salt (valid per RFC 5869 [5], treated as zeros); per-credential domain separation comes from the unique PRF salt, which yields unique input keying material. The HKDF `info` label is the frozen constant

$$\ell_{v1} = \text{"webauthn-prf-zktv vault key wrap v1"}.$$

A legacy label, `"TrustVault Vault Key Wrapping v1"`, survives only inside `fromTrustVaultRecord`, the one-shot migration adapter: it unwraps a TrustVault-PWA record under the legacy label and re-wraps the same vault key under ℓ_{v1} using the same PRF output, so migration needs no additional ceremony. Both labels are frozen: changing either would silently invalidate every record wrapped under it, so format evolution happens by introducing a new scheme identifier and label (e.g. a future `prf-v2`), never by mutating `v1`.

A password fallback scheme `pw-v1` covers clients where PRF is not viable (notably Android WebView, where PRF does not reach the Credential Manager): the wrap key is derived with memory-hard scrypt ($N=131072$, $r=8$, $p=1$ as a security floor) instead of HKDF-from-PRF; all other steps are identical.

4.2 Pseudo-Code

Algorithm 1 reflects `wrapSecret` in `src/core/wrap.ts` composed with the PRF ceremony (`enrollVault` in `src/webauthn/vault.ts`):

Algorithm 1 Vault key wrapping (`enrollVault` $\rightarrow$ `wrapSecret`, scheme `prf-v1`)

Require: WebAuthn credential C with PRF enabled, PRF salt s (32 random bytes), vault key K_v (32 bytes)

Ensure: Stored record R that does not reveal K_v without C

- 1: $prf \leftarrow \text{EvaluatePRF}(C, s) \triangleright 32$ bytes; `PrfResultMissingError` if absent or wrong length
 - 2: $K_w \leftarrow \text{HKDF-SHA256}(\text{ikm}=prf, \text{salt}=\varepsilon, \text{info}=\ell_{v1})$
 - 3: import K_w as non-extractable AES-256-GCM `CryptoKey`
 - 4: $\text{nonce} \leftarrow \text{Random}(12)$
 - 5: $\text{ciphertext} \leftarrow \text{AES-256-GCM-Encrypt}(K_w, \text{nonce}, K_v) \triangleright$ auth tag included
 - 6: zeroize $prf \triangleright$ **finally** block; K_w is non-extractable, never in JS memory
 - 7: $R \leftarrow (\text{scheme}=\text{prf-v1}, \text{ciphertext}, \text{nonce}, \text{salt}=s)$
 - 8: **return** R
-

A deliberate design point concerns zeroization scope: `wrapSecret` does *not* zeroize the caller's secret buffer. Ownership of K_v remains with the caller, which may legitimately need the raw key again — most commonly to wrap the same vault key a second time under `pw-v1` so that a password fallback record can sit alongside the PRF record. Zeroization inside the library is therefore scoped to the buffers the library itself creates: PRF outputs after key derivation, and the transient plaintext produced during unwrapping.

Algorithm 2 reflects `unlockVault` $\rightarrow$ `unwrapSecret`; note the replay verification step absent from generic treatments:

Algorithm 2 Vault key unwrapping (`unlockVault` $\rightarrow$ `unwrapSecret`)

Require: Credential id, stored record $R = (\text{prf-v1}, \text{ciphertext}, \text{nonce}, s)$, last stored signature counter c

Ensure: Non-extractable session `CryptoKey` K_v , new counter c'

- 1: $(prf, c') \leftarrow \text{EvaluatePRF}(C, s) \triangleright$ assertion ceremony with fresh 32-byte challenge
 - 2: verify `clientData: type=webauthn.get`, challenge and origin match; verify $c' > c$ (unless $c'=0$); else throw `ReplayError`
 - 3: $K_w \leftarrow \text{HKDF-SHA256}(\text{ikm}=prf, \text{salt}=\varepsilon, \text{info}=\ell_{v1})$
 - 4: $b \leftarrow \text{AES-256-GCM-Decrypt}(K_w, \text{nonce}, \text{ciphertext})$; on any failure throw `DecryptError` $\triangleright$ never distinguishes wrong key from corrupt data
 - 5: reject unless $|b| = 32$; import b as non-extractable AES-256-GCM `CryptoKey` K_v
 - 6: zeroize b and $prf \triangleright$ both in **finally** blocks
 - 7: **return** $(K_v, c') \triangleright$ caller persists c' for the next unlock
-

Failure paths are typed. A wrong key or tampered ciphertext raises `DecryptError`, which is intentionally generic and propagates no cause, so an observer cannot distinguish a wrong-key attempt from corrupted data. Malformed or scheme-mismatched records raise `RecordFormatError`, and challenge, origin, or counter anomalies raise `ReplayError`.

4.3 Zero-Knowledge Argument

The claim is that the complete persisted state — the record R and any vault ciphertexts encrypted under K_v — is computationally useless without a ceremony on the enrolled authenticator. The argument rests on three standard assumptions: the authenticator’s `hmac-secret` PRF is pseudo-random to parties without the credential’s internal key; HKDF-SHA256 is a secure key-derivation function [5]; and AES-256-GCM provides IND-CPA confidentiality and INT-CTXT integrity.

For `prf-v1` records, R contains only public values: the scheme tag, the AES-GCM ciphertext and nonce, and the PRF salt s . The salt is an *input* to the PRF, not a secret; by PRF security, the output at s is indistinguishable from random to anyone who cannot query the authenticator. Hence the wrap key $K_w = \text{HKDF}(\text{prf}, \varepsilon, \ell_{v1})$ is indistinguishable from a uniformly random AES-256 key, and by IND-CPA security the ciphertext reveals nothing about K_v . No other function of the PRF output or of K_v is ever persisted (Sect. 5), so stored-data compromise reduces to attacking AES-256-GCM with an unknown random key.

For `pw-v1` records the guarantee is necessarily weaker: the record is exactly as strong as the password under offline guessing, which is why the script cost parameters stored in `kdfParams` are treated as a floor, not a tunable. The domain-separation label ℓ_{v1} ensures that even if the same PRF output were ever used with another HKDF-based scheme, the derived keys would be independent.

4.4 Design Rationale: Deliberate Omissions and the Frozen v1 Format

Two natural hardening measures are deliberately absent from the v1 record format, and both omissions are governed by the same constraint: v1 is frozen, because any change to the key-derivation inputs or the ciphertext computation silently invalidates every record already persisted by deployed applications.

HKDF salt. RFC 5869 [5] recommends a salt where one is available, and the PRF salt s is a natural candidate for the HKDF salt field. v1 instead extracts with an empty salt and relies on s for per-credential separation *via the IKM*: distinct salts yield independent PRF outputs, so distinct credentials already derive independent wrap keys, and the versioned info label provides domain separation across schemes. The security loss is confined to extractor robustness against adversarially structured IKM — a non-concern here, since the IKM is the output of a PRF evaluated inside the authenticator. A future `prf-v2` scheme will nevertheless move s into the HKDF salt field to align with common practice, at the cost of a re-wrap migration.

AEAD context binding. `v1` passes no additional authenticated data (AAD) to AES-GCM. GCM’s authentication tag already rejects any modification of nonce or ciphertext; what AAD would add is *context* binding — preventing a valid record from being replayed in a different slot (e.g. swapping a `pw-v1` record into a `prf-v1` position, or moving a record between vaults). `v1` compensates structurally: `parseRecord` enforces scheme-specific invariants (`kdfParams` present iff `pw-v1`), unwrapping asserts the record’s scheme against the supplied key material before touching the cipher, and the `IndexedDB` composite key `[vaultId, scheme]` fixes each record’s position. A context swap therefore fails before or at decryption, but the binding is enforced by the validation layer rather than the AEAD. `prf-v2` is reserved to bind (version, scheme, salt, vaultId) as AAD, converting this structural argument into a cryptographic one.

4.5 Application-Layer Guidance Against Non-Cryptographic Leakage

The zero-knowledge argument of Sect. 4.3 covers only the persisted record and vault ciphertexts; it says nothing about surfaces above the storage layer. Fábrega et al. [4] show that E2EE password managers leak through exactly these surfaces under adversarial data injection. Applications adopting this library should additionally:

- **Never compress or deduplicate before encryption.** Compression ratios and dedup hits both leak information about plaintext content and must happen, if at all, only after AES-GCM encryption (where they are ineffective against ciphertext, which is the point).
- **Pad or bucket record sizes.** Ciphertext length is a direct function of plaintext length; secrets with a wide size distribution (e.g. free-text notes vs. fixed-width passwords) should be padded to fixed buckets before wrapping so stored record size does not correlate with secret content.
- **Keep telemetry and analytics out of the unlock/wrap path.** No metrics, crash reports, or usage analytics should be scoped to `enrollVault/unlockVault` calls or fire conditionally on their success/failure in a way that leaks timing or outcome to a third party.
- **Do not let vault content drive external fetches.** Icon, favicon, or meta-data lookups (a known injection vector) must never be triggered by, or scoped to, individual vault entries.
- **Treat IndexedDB write timing as an observable channel.** Batching or jittering writes avoids letting local instrumentation (a malicious extension, a compromised service worker) infer which record changed from write timing alone.

These are library-consumer obligations, not properties the library can enforce from inside `ZktvDb`, since padding and telemetry policy are application-specific.

5 IndexedDB Storage Architecture

5.1 Schema Overview

The library’s `ZktvDb` (raw IndexedDB with no wrapper dependency; database name `zktv`, version 1) creates three object stores:¹

- **vaults**: composite primary key [`vaultId`, `scheme`]. Each row holds `vaultId`, `scheme`, the JSON-serialized `WrappedSecretRecord` (version-tagged `v:1`, binary fields base64url-encoded), and `updatedAt`. The composite key lets one vault carry a `prf-v1` record and a `pw-v1` fallback record side by side; `loadWrappedVault` tries `prf-v1` first.
- **credentials**: primary key `credentialId` (base64url), with a secondary index on `vaultId`. Rows hold non-secret metadata only: the owning vault id, the last verified signature counter, the base64url PRF salt, the authenticator transports, and a creation timestamp.
- **meta**: key–value store (`keyPath`: `key`) reserved for application metadata and future migrations.

Every stored row round-trips through `serializeRecord/parseRecord`, so structural invariants (12-byte nonce, ≥ 16 -byte salt, ciphertext longer than the 16-byte GCM tag, `kdfParams` present iff `pw-v1`) are re-validated on every load and violations surface as `RecordFormatError`.

5.2 Schema Diagram

Figure 1 depicts the three stores and their relationship.

5.3 Migrations and Security Properties

`openVaultDb` accepts an `onUpgrade(db, oldVersion, tx)` hook that runs inside `onupgradeneeded` *after* the built-in stores are ensured, letting applications version their own stores without touching the core schema. The migration invariant is that no upgrade may introduce stored recomputable key inputs: only `WrappedSecretRecord` shapes and non-secret credential metadata are ever persisted. Lifecycle operations preserve this: `clearVault` deletes both scheme rows for a vault and every credential row bound to it via the `vaultId` index, and `securityWipe` clears all three stores (mirroring TrustVault’s security-wipe semantics). Key hygiene throughout: all derived and imported keys are non-extractable `CryptoKeys`, transient buffers are zeroized in `finally` blocks, and the signature counter persisted in `credentials` is updated after each verified assertion (`updateCounter`) to feed the replay check of Algorithm 2.

¹ The TrustVault-PWA *application* persists further stores at its own layer — notably breach-detection (HIBP) hash prefixes and ephemeral session metadata [11]. Those are application concerns outside the library schema described here, which ships exactly the three stores listed.




Fig. 1. IndexedDB schema created by `openVaultDb` (database `zktv`, version 1; `src/indexeddb/db.ts`). One vault may hold up to two wrapped records (`prf-v1` and the `pw-v1` fallback) via the composite primary key, and any number of PRF credentials via the `vaultId` index. Only wrapped records and non-secret credential metadata are ever persisted.

6 Implementation in `webauthn-prf-zktv`

The library [12] is strict-mode TypeScript, ESM-only, with `@noble/hashes` (`scrypt`) as its sole runtime dependency; all other cryptography uses the Web Crypto API. It ships three sub-path entry points:

- `webauthn-prf-zktv` (core; Node ≥ 20 and browsers): the wrap and unwrap functions of Sect. 4, HKDF and `scrypt` derivation, record (de)serialization with structural validation, the TrustVault migration adapter, and a typed error hierarchy (`ZktvError`) whose codes distinguish PRF unavailability, cancelled ceremonies, missing PRF results, replay anomalies, decryption failure, record-format violations, and storage errors.
- `webauthn-prf-zktv/webauthn` (browsers only): capability detection and ceremonies. Detection is deliberately non-invasive. `detectPrfSupport` is tri-state (supported, unsupported, unknown) via `getClientCapabilities()` and never creates throwaway credentials to probe; the stricter viability check, `isPrfViableOnThisClient`, additionally rejects Android WebView — where PRF does not reach the Credential Manager — and requires a user-verifying platform authenticator. Enrollment is *adaptive*: `enrollPrfCredential` requests `prf.eval` at creation time, so clients that evaluate PRF during registration (e.g. recent Chrome with Windows Hello) enroll in a single ceremony; otherwise it hard-verifies `prf.enabled` in the extension outputs, aborting cleanly if absent, and runs one assertion ceremony to obtain

the PRF output. `evaluatePrf` enforces the 32-byte PRF result length and performs the replay verification of Algorithm 2; `enrollVault` and `unlockVault` compose ceremony, derivation, and (un)wrapping into single calls that zeroize the PRF output in `finally` blocks.

- `webauthn-prf-zktv/indexeddb` (browsers only): the `ZktvDb` storage layer of Sect. 5 (Fig. 1).

The public API is exercised by unit tests covering the failure paths of Sect. 4 (wrong key, malformed records, replayed assertions) alongside the success paths. The core wrap entry point, verbatim from `src/core/wrap.ts`:

Listing 1.1. `wrapSecret`: the three wrap paths (PRF, password, caller-supplied wrap key).

```
export async function wrapSecret(options: WrapOptions): Promise<
  WrappedSecretRecord> {
  if ('prfOutput' in options) {
    const wrapKey = await deriveWrapKeyFromPrf(options.prfOutput);
    return encryptRecord(wrapKey, 'prf-v1', options.prfSalt, options.
      secret);
  }
  if ('password' in options) {
    const salt = options.salt ?? generateSalt();
    const kdfParams = options.kdfParams ?? DEFAULT_SCRIPT_PARAMS;
    const wrapKey = await deriveWrapKeyFromPassword(options.password,
      salt, kdfParams);
    const record = await encryptRecord(wrapKey, 'pw-v1', salt, options.
      secret);
    return { ...record, kdfParams };
  }
  return encryptRecord(options.wrapKey, options.scheme, options.salt,
    options.secret);
}
```

Listing 1.2. HKDF derivation with the frozen v1 domain-separation label (`src/core/derive.ts`).

```
export const HKDF_INFO_V1 = new TextEncoder().encode(
  'webauthn-prf-zktv vault key wrap v1',
);

export async function deriveWrapKeyFromPrf(prfOutput, info = HKDF_INFO_V1) {
  {
    const ikm = await crypto.subtle.importKey('raw', prfOutput, 'HKDF',
      false, ['deriveKey']);
    return crypto.subtle.deriveKey(
      { name: 'HKDF', hash: 'SHA-256', salt: new Uint8Array(0), info },
      ikm,
      { name: 'AES-GCM', length: 256 },
      false, // non-extractable
    );
  }
}
```

```

    ['encrypt', 'decrypt'],
  );
}

```

7 Post-Quantum Considerations

The storage layer of this architecture is built entirely from symmetric primitives — HMAC-SHA-256 (inside the authenticator PRF), HKDF-SHA256, scrypt, and AES-256-GCM. Against a quantum adversary, Grover-type speedups at most halve the effective security exponent, leaving AES-256 and SHA-256-based constructions with comfortable margins; NIST’s transition guidance accordingly treats AES-256 and SHA-256/384 as quantum-adequate [9], a position echoed by recent post-quantum migration surveys [2]. Harvest-now, decrypt-later attacks against a captured IndexedDB dump are therefore not materially assisted by future quantum capability.

The asymmetric layer is different: today’s passkeys sign with elliptic-curve or RSA algorithms (the library requests ES256 and RS256), which fall to a cryptographically relevant quantum computer. However, that layer protects *authentication*, not the confidentiality of stored records: breaking the signature scheme lets an attacker forge assertions toward a relying party, but does not reveal the `hmac-secret` key inside the authenticator, and hence not the PRF output or K_v . Because the wrapping scheme is agnostic to the credential’s signature algorithm, PQ-capable passkeys (e.g. ML-DSA-based, as they appear in authenticators) can be adopted without any change to the record format; a future `prf-v2` scheme identifier is reserved for the case where the KDF or cipher itself must move.

8 Evaluation and Discussion

The reference implementation is evaluated along three axes. *Security*: the zero-knowledge argument of Sect. 4.3 covers the persisted state; the unit-test suite additionally pins the negative behaviors the argument relies on — wrong-key decryption is indistinguishable from corruption, malformed records never reach the cipher, and non-increasing signature counters abort the unlock. *Performance*: unlock cost is dominated by the WebAuthn ceremony itself (user verification takes seconds), against which one HKDF derivation and one AES-GCM decryption of a 32-byte payload are negligible. The password fallback is deliberately expensive: scrypt with $N=2^{17}=131,072$, $r=8$, $p=1$ requires a $128 \cdot N \cdot r = 128$ MiB working set and costs a few hundred milliseconds on desktop-class hardware, but can reach multi-second latency and cause memory pressure on low-end mobile devices and WebViews. These parameters are a security floor, not a tunable: the wrap KDF is precisely what an offline attacker attacks, so adaptive parameter selection is admissible only *upward*; applications targeting constrained devices should prefer the PRF path, whose cost is independent of the KDF floor. *Usability*: PRF-based unlock replaces master-password entry with a platform biometric

gesture, but availability is uneven across browsers, platforms, and embedded WebViews; the tri-state capability detection and the `pw-v1` fallback exist precisely so that applications degrade gracefully rather than probe destructively. *Failure semantics*: record and counter writes are single IndexedDB transactions, so a crash mid-operation leaves either the old or the new row, never a torn one; a lost counter update leaves a stale (lower) stored counter, which can only widen the acceptance window of the next unlock by one assertion — it can never render the vault un-unlockable, so replay protection degrades toward availability rather than lockout. *Scope of the zero-knowledge claim*: the guarantee of Sect. 4.3 covers persisted state only; system-level channels demonstrated against E2EE password managers — telemetry, caching, and pre-encryption compression or deduplication [4] — as well as record-size and timing patterns in IndexedDB remain application responsibilities (Sect. 4.5 lists the corresponding app-layer rules). A quantitative cross-device benchmark of ceremony latency, PRF availability, and script cost on constrained devices, together with a game-based formalization of the wrapping scheme, is left as future work; interoperability test vectors for both record schemes are published with the library to let independent implementations validate compatibility.

9 Conclusion

This paper formalized a WebAuthn PRF-based vault key wrapping scheme — PRF output as HKDF input keying material under a frozen, versioned domain-separation label, wrapping a non-extractable AES-256-GCM vault key — together with an IndexedDB storage architecture whose entire persisted state provably contains no recomputable key inputs. The scheme integrates client-side replay verification, a memory-hard password fallback for clients where PRF is unavailable, and a ceremony-free migration path from the predecessor TrustVault-PWA format. The open-source `webauthn-prf-zktv` library [12] packages these patterns behind small, typed entry points so that other web applications can adopt PRF-based zero-knowledge vaults without re-deriving the storage and key-hygiene discipline described here.

References

1. holstorage: client-side file encryption with the WebAuthn PRF extension. Open-source project, 2024. Available at <https://github.com/holstorage/webauthn-prf>.
2. Yaser Baseri, Abdelhakim Hafid, and Arash Habibi Lashkari. Future-proofing cloud security against quantum attacks: Risk, transition, and mitigation strategies. arXiv:2509.15653, 2025.
3. Bitwarden. PRF WebAuthn and its role in passkeys. Bitwarden Blog, 2024.
4. Andrés Fábrega, Armin Namavari, Rachit Agarwal, Ben Nassi, and Thomas Ristenpart. Exploiting leakage in password managers via injection attacks. arXiv:2408.07054, 2024.

5. Hugo Krawczyk and Pasi Eronen. HMAC-based extract-and-expand key derivation function (HKDF). RFC 5869, IETF, 2010.
6. Klaudia Krawiecka, Arseny Kurnikov, Andrew Paverd, Mohammad Mannan, and N. Asokan. SafeKeeper: Protecting web passwords using trusted execution environments. arXiv:1709.01261, 2017.
7. Hannah Li and David Evans. Horcrux: A password manager for paranoids. arXiv:1706.05085, 2017.
8. Matthew Miller. The WebAuthn PRF extension. SimpleWebAuthn documentation, 2024.
9. Dustin Moody et al. Transition to post-quantum cryptography standards. NIST Internal Report 8547 (initial public draft), National Institute of Standards and Technology, 2024.
10. Vivek Nair and Dawn Song. MFDPG: Multi-factor authenticated password management with zero stored secrets. arXiv:2306.14746, 2023.
11. Ian Pinto. TrustVault-PWA: security architecture and threat model documentation. GitHub repository, 2026. Available at <https://github.com/opnsrcntrbtr/TrustVault-PWA>.
12. Ian Pinto. webauthn-prf-zktv: WebAuthn PRF-backed vault key wrapping with zero-knowledge IndexedDB storage, v0.1.0. npm package and GitHub repository, 2026. Reference implementation for this paper; also published as **webauthn-prf-zktv** on npm.
13. Colin Roberts, Vivek Nair, and Dawn Song. Wrangling entropy: Next-generation multi-factor key derivation, credential hashing, and credential generation functions. arXiv:2509.05893, 2025.
14. W3C Web Authentication Working Group. Web authentication: An API for accessing public key credentials, level 3. W3C Working Draft, 2025. §10.1.4 defines the **prf** client extension.
15. Yubico. The WebAuthn PRF extension. Yubico Developers documentation, 2024.
16. Clemens Zeidler and Muhammad Rizwan Asghar. AuthStore: Password-based authentication and encrypted data storage in untrusted environments. arXiv:1805.05033, 2018.